

CS2010: Data Structures and Algorithms II

Substring Search – part 2

Ivana.Dusparic@scss.tcd.ie

Substring search algorithms – part 2

- › Boyer-Moore
- › Rabin-Karp
- › Suffix arrays/suffix trees and LCPs

So far

- › Brute force
 - $M \times N$ performance
 - Back up
- › KMP
 - $2N$ - linear
 - No backup
 - Extra space – $M \times R$

(N length of a string, M length of a substring we're search for, R radix)

- › can we do better?



Boyer-Moore



Boyer-Moore

- › Big idea – when find a character not in the pattern, can skip up to M characters (so no need to loop through all N characters)
 - Mismatched character heuristic
 - Don't look through characters in order, start from the back and look at the last character in the pattern first and see if it's a match, or in the pattern at all
- › Uses backup

Boyer-Moore

Intuition.

- Scan characters in pattern from right to left.
- Can skip as many as M text chars when finding one not in the pattern.

i	j	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23
	<i>text</i> →	F	I	N	D	I	N	A	H	A	Y	S	T	A	C	K	N	E	E	D	L	E	I	N	A
0	5	N	E	E	D	L	E																		
5	5						N	E	E	D	L	E													
11	4												N	E	E	D	L	E							
15	0																N	E	E	D	L	E			


← *pattern*

↖ *return i = 15*

Boyer-Moore: mismatched character heuristic

Q. How much to skip?

Case 1. Mismatch character not in pattern.



mismatch character 'T' not in pattern: increment i one character beyond 'T'

Boyer-Moore: mismatched character heuristic

Q. How much to skip?

Case 2a. Mismatch character in pattern.

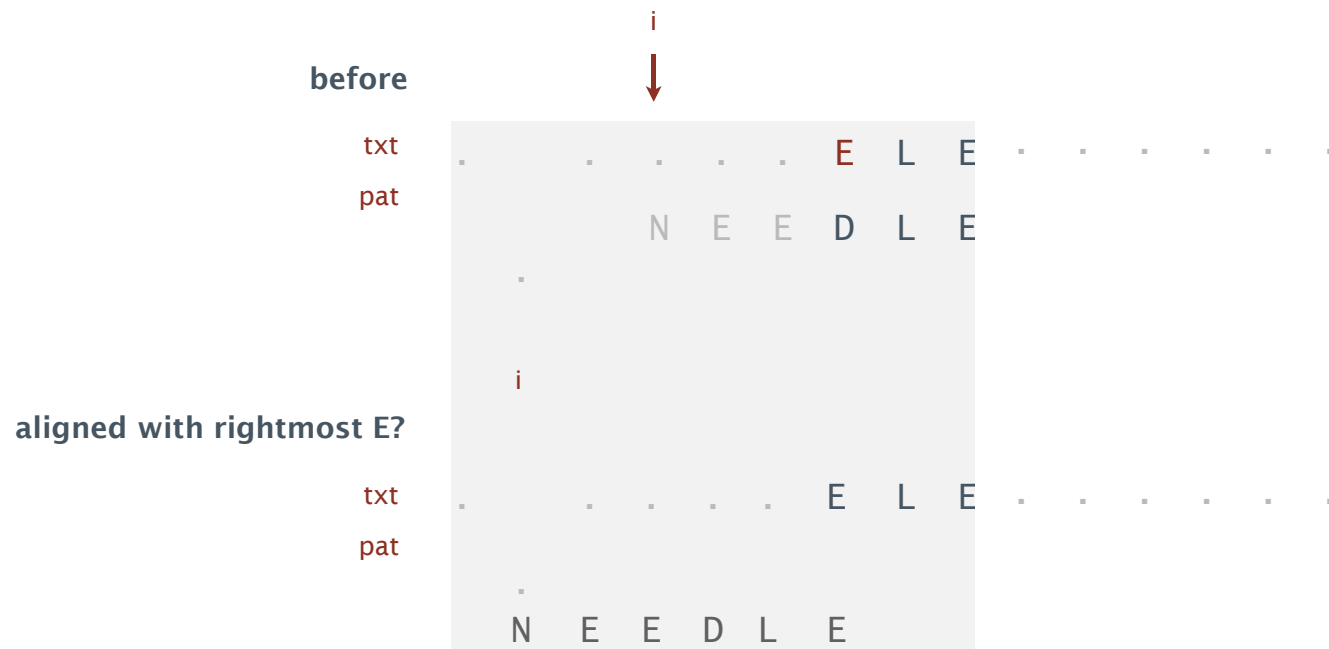


mismatch character 'N' in pattern: align text 'N' with rightmost pattern 'N'

Boyer-Moore: mismatched character heuristic

Q. How much to skip?

Case 2b. Mismatch character in pattern (but heuristic no help).




mismatch character 'E' in pattern: align text 'E' with rightmost pattern 'E' ?

Boyer-Moore: mismatched character heuristic

Q. How much to skip?

Case 2b. Mismatch character in pattern (but heuristic no help).



mismatch character 'E' in pattern: increment i by 1

Boyer-Moore: mismatched character heuristic

Q. How much to skip?

A. Precompute index of rightmost occurrence of character c in pattern.
(-1 if character not in pattern)

```
right = new int[R];  
for (int c = 0; c < R; c++)  
    right[c] = -1;  
for (int j = 0; j < M; j++)  
    right[pat.charAt(j)] = j;
```

		N	E	E	D	L	E	
	c	0	1	2	3	4	5	right[c]
A		-1	-1	-1	-1	-1	-1	-1
B		-1	-1	-1	-1	-1	-1	-1
C		-1	-1	-1	-1	-1	-1	-1
D		-1	-1	-1	3	3	3	3
E		-1	-1	1	2	2	5)	5
...								-1
L		-1	-1	-1	-1	4	4	4
M		-1	-1	-1	-1	-1	-1	-1
N		-1	0	0	0	0	0	0
...								-1

Boyer-Moore skip table computation

Precomputing the index of right-most occurrence

```
// position of rightmost occurrence of c in the pattern  
right = new int[R];  
for (int c = 0; c < R; c++)  
    right[c] = -1;  
for (int j = 0; j < pattern.length; j++)  
    right[pattern[j]] = j;
```

Boyer-Moore: Java implementation


```
public int search(String txt)
{
    int N = txt.length();
    int M = pat.length();
    int skip;
    for (int i = 0; i <= N-M; i += skip)
    {
        skip = 0;
        for (int j = M-1; j >= 0; j--)
        {
            if (pat.charAt(j) != txt.charAt(i+j))
            {
                skip = Math.max(1, j - right[txt.charAt(i+j)]);
                break;
            }
        }
        if (skip == 0) return i;
    }
    return N;
}
```

compute skip value

in case other term is nonpositive

match

Boyer-Moore: analysis

Property. Substring search with the Boyer-Moore mismatched character heuristic takes about $\sim N/M$ character compares to search for a pattern of length M in a text of length N .  sublinear!

What's the worst case input/performance of Boyle-Moore?

Turning point.

Boyer-Moore: analysis

Worst-case. Can be as bad as $\sim MN$.

i	skip		0	1	2	3	4	5	6	7	8	9
		txt	B	B	B	B	B	B	B	B	B	B
0	0		A	B	B	B	B		pat			
1	1			A	B	B	B	B				
2	1				A	B	B	B	B			
3	1					A	B	B	B	B		
4	1						A	B	B	B	B	
5	1							A	B	B	B	B



Rabin-Karp



Rabin-Karp

Basic idea = modular hashing.

- Compute a hash of `pat[0..M-1]`.
- For each `i`, compute a hash of `txt[i..M+i-1]`.
- If pattern hash = text substring hash, check for a match.

pat.charAt(i)																										
i	0	1	2	3	4																					
	2	6	5	3	5	% 997 = 613																				
						txt.charAt(i)																				
i	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15										
	3	1	4	1	5	9	2	6	5	3	5	8	9	7	9	3										
0	3	1	4	1	5	% 997 = 508																				
1		1	4	1	5	9	% 997 = 201																			
2			4	1	5	9	2	% 997 = 715																		
3				1	5	9	2	6	% 997 = 971																	
4					5	9	2	6	5	% 997 = 442																
5						9	2	6	5	3	% 997 = 929															
6	← return i = 6						2	6	5	3	5	% 997 = 613														

match

modular hashing with $R = 10$ and $\text{hash}(s) = s \pmod{997}$

Modular hashing

- › Remember hash tables
- › *hash function* - map data of arbitrary size to data of fixed size
- › Eg map any value to an index in the array
- › With **modular hashing**, the hash function is simply $H(k) = k \bmod m$ for some m . The value k is an integer hash code generated from the key (generally used with positive integers)

```
int h(int k, int M) {  
    return k % M;  
}
```

Rabin-Karp – modular hashing

- › Won't it overflow, for large search substrings?

Math trick. To keep numbers small, take intermediate results modulo Q .

Ex.
$$\begin{aligned}(10000 + 535) * 1000 & \pmod{997} \\ &= (30 + 535) * 3 \pmod{997} \\ &= 1695 \pmod{997} \\ &= 698 \pmod{997}\end{aligned}$$

$$(a + b) \bmod Q = ((a \bmod Q) + (b \bmod Q)) \bmod Q$$

$$(a * b) \bmod Q = ((a \bmod Q) * (b \bmod Q)) \bmod Q$$

two useful modular arithmetic identities

Efficiently computing the hash function

Modular hash function. Using the notation t_i for `txt.charAt(i)`, we wish to compute

$$x_i = t_i R^{M-1} + t_{i+1} R^{M-2} + \dots + t_{i+M-1} R^0 \pmod{Q}$$

Intuition. M -digit, base- R integer, modulo Q .

Horner's method. Linear-time method to evaluate degree- M polynomial.

	<code>pat.charAt()</code>				
<code>i</code>	0	1	2	3	4
	2	6	5	3	5
0	2	% 997 = 2			
1	2	6	% 997 = (2*10 + 6) % 997 = 26		
2	2	6	5	% 997 = (26*10 + 5) % 997 = 265	
3	2	6	5	3	% 997 = (265*10 + 3) % 997 = 659
4	2	6	5	3	5 % 997 = (659*10 + 5) % 997 = 613

```
// Compute hash for M-digit key
private long hash(String key, int M)
{
    long h = 0;
    for (int j = 0; j < M; j++)
        h = (h * R + key.charAt(j)) % Q;
    return h;
}
```

$$\begin{aligned} 26535 &= 2*10000 + 6*1000 + 5*100 + 3*10 + 5 \\ &= (((2) * 10 + 6) * 10 + 5) * 10 + 3) * 10 + 5 \end{aligned}$$

Efficiently computing the hash function

Challenge. How to efficiently compute x_{i+1} given that we know x_i .

$$x_i = t_i R^{M-1} + t_{i+1} R^{M-2} + \dots + t_{i+M-1} R^0$$

$$x_{i+1} = t_{i+1} R^{M-1} + t_{i+2} R^{M-2} + \dots + t_{i+M} R^0$$

Key property. Can update "rolling" hash function in constant time!

$$x_{i+1} = \underbrace{(x_i)}_{\text{current value}} - \underbrace{t_i R^{M-1}}_{\text{subtract leading digit}} \underbrace{R}_{\text{multiply by radix}} + \underbrace{t_{i+M}}_{\text{add new trailing digit}} \quad (\text{can precompute } R^{M-1})$$

i	...	2	3	4	5	6	7	...
current value	1	4	1	5	9	2	6	5
new value		4	1	5	9	2	6	5
		4	1	5	9	2	text	
		4	1	5	9	2		
		4	1	5	9	2	current value	
-		4	0	0	0	0	subtract leading digit	
			1	5	9	2	multiply by radix	
				*	1	0		
		1	5	9	2	0	add new trailing digit	
					+	6		
		1	5	9	2	6	new value	

Rabin-Karp substring search example

First R entries: Use Horner's rule.

Remaining entries: Use rolling hash (and % to avoid overflow).

i	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
	3	1	4	1	5	9	2	6	5	3	5	8	7	9	3	
0	3	% 997 = 3										9				
1	3	1	% 997 = (3*10 + 1) % 997 = 31													
2	3	1	4	% 997 = (31*10 + 4) % 997 = 314												
3	3	1	4	1	% 997 = (314*10 + 1) % 997 = 150											
4	3	1	4	1	5	% 997 = (150*10 + 5) % 997 = 508										
5		1	4	1	5	9	% 997 = ((508 + 3*(997 - 30))*10 + 9) % 997 = 201									
6			4	1	5	9	2	% 997 = ((201 + 1*(997 - 30))*10 + 2) % 997 = 715								
7				1	5	9	2	6	% 997 = ((715 + 4*(997 - 30))*10 + 6) % 997 = 971							
8					5	9	2	6	5	% 997 = ((971 + 1*(997 - 30))*10 + 5) % 997 = 442						
9						9	2	6	5	3	% 997 = ((442 + 5*(997 - 30))*10 + 3) % 997 = 929					
10							2	6	5	3	5	% 997 = ((929 + 9*(997 - 30))*10 + 5) % 997 = 613				

$\leftarrow \text{return } i - M + 1 = 6$

-30 (mod 997) = 997 - 30 10000 (mod 997) = 30

Horner's rule
 rolling hash
 match

Rabin-Karp: Java implementation

```
public class RabinKarp
{
    private long patHash;    // pattern hash
                           // value
    private int M;          // pattern length
    private long Q;         // modulus
    private int R;          // radix
    private long RM1;       //  $R^{M-1} \% Q$ 
    public RabinKarp(String pat) {
        M = pat.length();
        R = 256;
        Q = longRandomPrime();

        RM1 = 1;
        for (int i = 1; i <= M-1; i++)
            RM1 = (R * RM1) % Q;
        patHash = hash(pat, M);
    }

    private long hash(String key, int M)
    { /* as before */ }

    public int search(String txt)
    { /* see next slide */ }
}
```

← a large prime
(but avoid overflow)


← precompute $R^{M-1} \pmod Q$

Rabin-Karp: Java implementation (continued)

Monte Carlo version. Return match if hash match.

```
public int search(String txt)
{
    int N = txt.length();
    int txtHash = hash(txt, M);
    if (patHash == txtHash) return 0;
    for (int i = M; i < N; i++)
    {
        txtHash = (txtHash + Q - RM*txt.charAt(i-M) % Q) % Q;
        txtHash = (txtHash*R + txt.charAt(i)) % Q;
        if (patHash == txtHash) return i - M + 1;
    }
    return N;
}
```

check for hash collision
using rolling hash function



Las Vegas version. Check for substring match if hash match;
continue search if false collision.

Rabin-Karp analysis

Monte Carlo version.

- Always runs in linear time.
- Extremely likely to return correct answer (but not always!).

Las Vegas version.

- Always returns correct answer.
- Extremely likely to run in linear time (but worst case is MN).

Rabin-Karp

Advantages.

- Extends to 2d patterns.
- Extends to finding multiple patterns.

Disadvantages.

- Arithmetic ops slower than char compares.
- Las Vegas version requires backup.
- Poor worst-case guarantee.

Substring search cost summary


Cost of searching for an M -character pattern in an N -character text.

algorithm	version	operation count		backup in input?	correct?	extra space
		guarantee	typical			
brute force	—	MN	$1.1N$	<i>yes</i>	<i>yes</i>	1
Knuth-Morris-Pratt	<i>full DFA (Algorithm 5.6)</i>	$2N$	$1.1N$	<i>no</i>	<i>yes</i>	MR
	<i>mismatch transitions only</i>	$3N$	$1.1N$	<i>no</i>	<i>yes</i>	M
Boyer-Moore	<i>full algorithm</i>	$3N$	N/M	<i>yes</i>	<i>yes</i>	R
	<i>mismatched char heuristic only (Algorithm 5.7)</i>	MN	N/M	<i>yes</i>	<i>yes</i>	R
Rabin-Karp [†]	<i>Monte Carlo (Algorithm 5.8)</i>	$7N$	$7N$	<i>no</i>	<i>yes[†]</i>	1
	<i>Las Vegas</i>	$7N^+$	$7N$	<i>yes</i>	<i>yes</i>	1

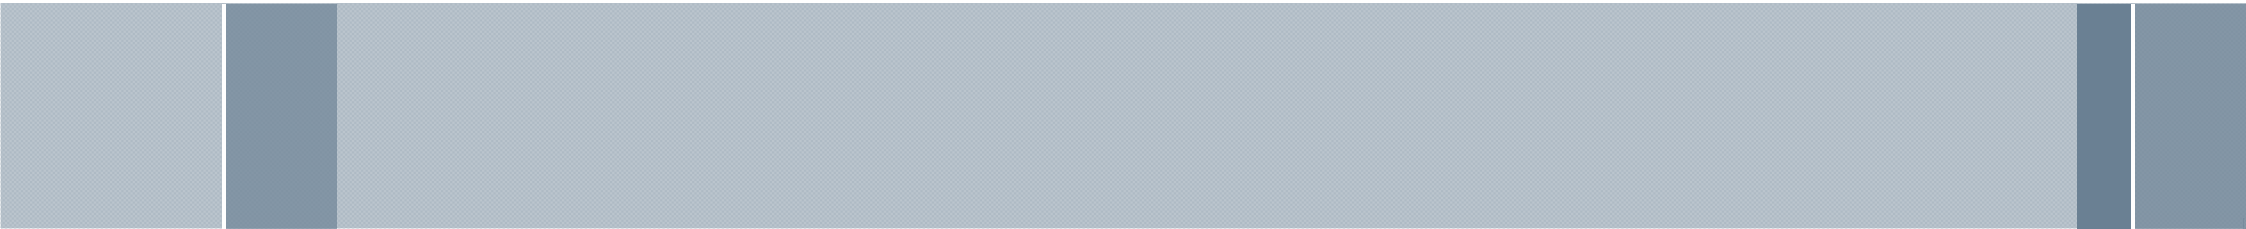
[†] probabilistic guarantee, with uniform hash function

So which algorithm should I use?

- › Java String.contains() method – brute force
 - Very compact, very little operations inside the loop, so loop runs fast. As there is no overhead it performs better when searching short strings.
- › Boyer-Moore
 - grep
 - Works well for long search patterns
 - The more distinct letters in the strings, the better the positive impact
 - backup
- › KMP
 - Small alphabet, repeated subpatterns
 - No backup



Suffix Arrays and LCP Arrays



Suffix Tree

- › A tree containing all the suffixes of the given text as their keys and positions in the text as their values
- › Space and time linear to length of the string
- › Speeds up operations such as: substring search, matches for regular expression patterns, longest common substring etc
- › Cost - storing a string's suffix tree typically requires significantly more space than storing the string itself.

- › slower but takes up less space and better cache performance


[illegible]

Sorted Suffix Array

Sorted Suffix										Index	
a										10	
a	b	r	a							7	
a	b	r	a	c	a	d	a	b	r	a	0
a	c	a	d	a	b	r	a				3
a	d	a	b	r	a						5
b	r	a									8
b	r	a	c	a	d	a	b	r	a		1
c	a	d	a	b	r	a					4
d	a	b	r	a							6
r	a										9
r	a	c	a	d	a	b	r	a			2

Longest Common Prefix Array

- › stores the lengths of the longest common prefixes (LCPs) between all pairs of consecutive suffixes in a sorted suffix array.



More on this and practice problems

› Robert Sedgewick and Kevin Wayne

<https://algs4.cs.princeton.edu/63suffix/>

› Suffix arrays – a programming contest approach

<http://web.stanford.edu/class/cs97si/suffix-array.pdf>